BOSSCODER ACADEMY

LEARN

MACHINE LEARNING

in 20 Days

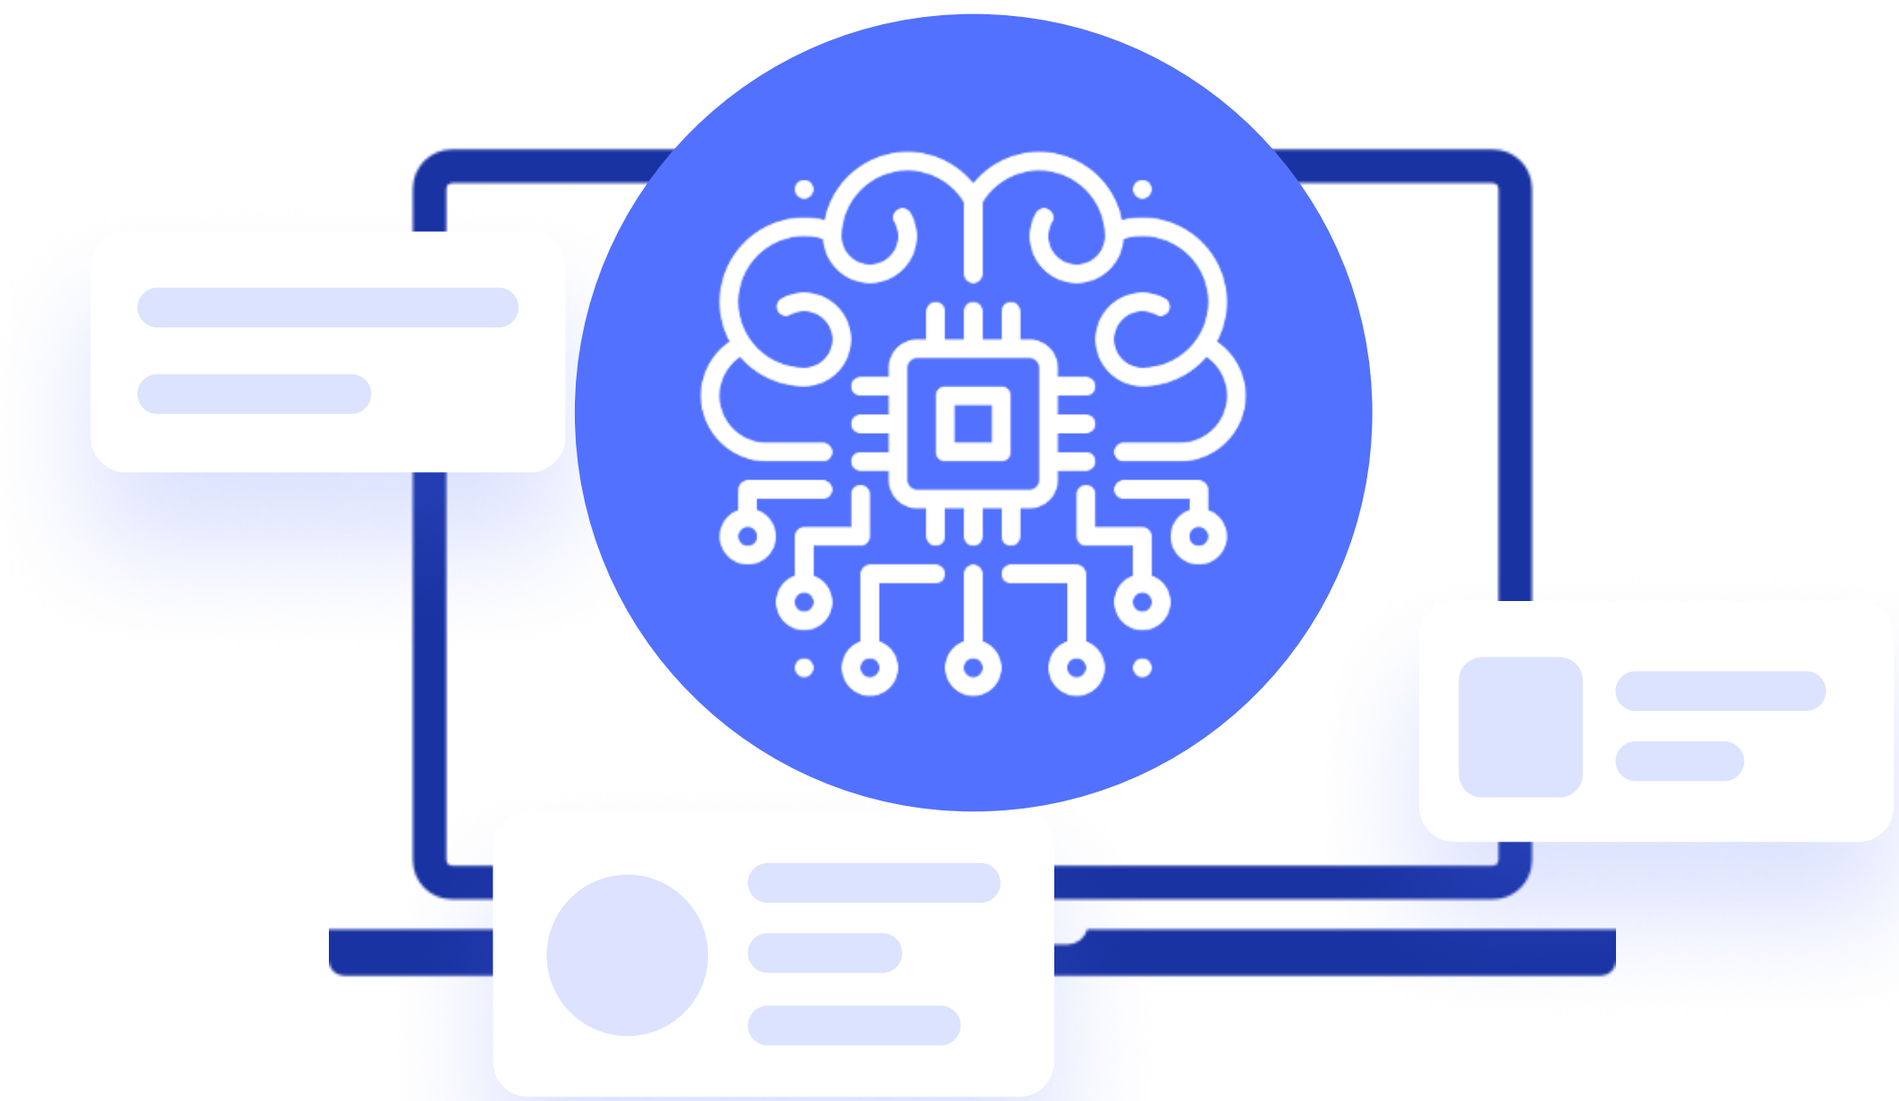

Disclaimer

This 20-day plan is designed to provide a strong foundation in Machine Learning.

Take the help of this doc and have a deep and practical understanding of key machine learning concepts and techniques.

DAY 1

Introduction to Machine Learning

Goals:

Understand what machine learning is and its various types.

Topics:

- Definition of machine learning
- Types of machine learning: supervised, unsupervised, and reinforcement learning

Resources:

- Machine Learning by Andrew Ng on Coursera
- Machine Learning Crash Course by Google

Practice Questions:

1. Define machine learning and explain how it differs from traditional programming.
2. List and describe the three main types of machine learning.

DAY 2

Tools and Software for Machine Learning

Goals:

Get familiar with the tools and programming languages used in machine learning.

Topics:

- Introduction to Python
- Overview of machine learning libraries (Scikit-learn, TensorFlow, PyTorch)

Resources:

- Python Machine Learning Tutorial for Beginners on Kaggle
- Scikit-learn Documentation

Practice Questions:

1. Why is Python favored for machine learning?
2. Compare TensorFlow and PyTorch. Which situations might favor one over the other?

DAY 3

Data Preprocessing

Goals:

Learn how to prepare data for machine learning models.

Topics:

- Handling missing data
- Data normalization and standardization
- Feature encoding

Resources:

- Data Preprocessing in Python on Real Python
- Scikit-learn Preprocessing Data

Practice Questions:

1. How do you handle missing values in a dataset using Pandas?
2. What is the difference between normalization and standardization?

DAY 4

Regression Analysis

Goals:

Understand and apply simple linear regression and multiple regression models.

Topics:

- Linear regression
- Multiple regression
- Model evaluation metrics (MSE, RMSE)

Resources:

- Linear Regression in Python on Stack Abuse
- Polynomial Regression on GeeksforGeeks

Practice Questions:

1. What assumptions must be met for linear regression to be effective?
2. How do you evaluate the performance of a regression model?

DAY 5

Classification Techniques

Goals:

Explore classification algorithms and their applications.

Topics:

- Logistic regression
- K-Nearest Neighbors (KNN)
- Support Vector Machines (SVM)

Resources:

- Decision Tree Classification on DataCamp
- Support Vector Machines on Scikit-learn

Practice Questions:

1. Compare and contrast logistic regression and SVM.
2. What are the main parameters that influence KNN performance?

DAY 6

Decision Trees and Random Forests

Goals:

Understand decision trees and how ensemble methods like random forests improve model performance.

Topics:

- Building decision trees
- Random forests
- Overfitting and model tuning

Resources:

- Random Forests in Python on Medium
- Decision Trees and Random Forests on Kaggle

Practice Questions:

1. Explain how a random forest model reduces the risk of overfitting.
2. What are the key hyperparameters in a random forest model?

DAY 7

Clustering Algorithms

Goals:

Learn about unsupervised learning techniques, focusing on clustering.

Topics:

- K-means clustering
- Hierarchical clustering
- DBSCAN

Resources:

- K-means Clustering in Python on Towards Data Science
- Hierarchical Clustering in Python on Scikit-learn

Practice Questions:

1. Describe a real-world application of clustering algorithms.
2. How do you determine the number of clusters in K-means?

DAY 8

Neural Networks Basics

Goals:

Gain a basic understanding of neural networks.

Topics:

- Anatomy of neural networks: neurons, layers, activation functions
- Forward propagation
- Loss functions

Resources:

- Neural Networks from Scratch on Sentdex
- A Beginner's Guide to Neural Networks on Medium

Practice Questions:

1. What role do activation functions play in a neural network?
2. Explain the concept of forward propagation.

DAY 9

Deep Learning with TensorFlow

Goals:

Explore deep learning frameworks and build a simple model with TensorFlow.

Topics:

- Introduction to TensorFlow
- Building a basic neural network in TensorFlow
- Understanding TensorFlow data pipelines

Resources:

- TensorFlow Official Tutorials
- Deep Learning with TensorFlow on Coursera

Practice Questions:

1. Build a TensorFlow model that classifies basic hand-written digits.
2. What is a tensor, and why is it fundamental to TensorFlow?

DAY 10

Convolutional Neural Networks (CNNs)

Goals:

Understand the specifics of CNNs and their applications in image processing.

Topics:

- Architecture of CNNs
- Pooling layers
- Convolution operations

Resources:

- Convolutional Neural Networks in Python on DataCamp
- CS231n: Convolutional Neural Networks for Visual Recognition on Stanford University

Practice Questions:

1. What are the advantages of using CNNs over traditional neural networks for image data?
2. Describe how pooling layers work in CNNs.

DAY 11

Recurrent Neural Networks (RNNs) and LSTMs

Goals:

Learn about RNNs and their use in sequence modeling.

Topics:

- Basics of RNNs
- Problems with RNNs (vanishing gradient)
- Long Short-Term Memory (LSTM) networks

Resources:

- Understanding LSTM Networks on colah's blog
- Recurrent Neural Networks by Example in Python on Towards Data Science

Practice Questions:

1. What is the vanishing gradient problem, and how do LSTMs address it?
2. Explain the differences between RNNs and LSTMs.

DAY 12

Natural Language Processing (NLP)

Goals:

Explore the basics of NLP and how machine learning is applied to text.

Topics:

- Text preprocessing techniques
- Bag of words and TF-IDF
- Basic sentiment analysis

Resources:

- Natural Language Processing with Python on NLTK
- A Comprehensive Guide to Text Preprocessing on Towards Data Science

Practice Questions:

1. How does TF-IDF work, and what does it measure?
2. Perform a simple sentiment analysis using a predefined library.

DAY 13

Natural Language Processing (NLP)

Goals:

Introduction to reinforcement learning and its key concepts.

Topics:

- Definition of reinforcement learning
- Key components: agents, environments, states, actions, rewards
- Q-learning

Resources:

- Reinforcement Learning: An Introduction on Sutton & Barto
- Deep Reinforcement Learning on FreeCodeCamp

Practice Questions:

1. Describe a scenario where reinforcement learning is applicable.
2. What is Q-learning, and how does it differ from other machine learning techniques?

DAY 14

Model Evaluation and Selection

Goals:

Learn how to evaluate and select machine learning models based on their performance.

Topics:

- Cross-validation
- ROC curves and AUC
- Confusion matrix

Resources:

- Model Evaluation Techniques in Machine Learning on Analytics Vidhya
- Scikit-learn Model Evaluation Documentation

Practice Questions:

1. What is cross-validation, and why is it important?
2. How does an ROC curve help in evaluating classification models?

DAY 15

Ensemble Methods

Goals:

Understand how ensemble methods combine predictions from different models to improve accuracy.

Topics:

- Bagging, boosting, and stacking
- AdaBoost and Gradient Boosting Machines (GBM)
- Model diversity and ensemble effectiveness

Resources:

- Ensemble Learning to Improve Machine Learning Results on StatQuest
- Introduction to Ensemble Learning on EliteDataScience

Practice Questions:

1. Explain how boosting is different from bagging.
2. What is stacking, and how can it improve model performance?

DAY 16

Feature Engineering and Dimensionality Reduction

Goals:

Learn techniques for feature extraction and reducing the dimensionality of data.

Topics:

- Feature engineering techniques
- Principal Component Analysis (PCA)
- t-SNE

Resources:

- Feature Engineering for Machine Learning on Udemy
- PCA Using Python (scikit-learn) on Towards Data Science

Practice Questions:

1. Describe how PCA reduces the dimensionality of data.
2. What are the benefits and drawbacks of using t-SNE over PCA?

DAY 17

Hyperparameter Tuning and Optimization

Goals:

Understand how to optimize machine learning models through hyperparameter tuning.

Topics:

- Grid search
- Random search
- Bayesian optimization

Resources:

- Hyperparameter Tuning the Random Forest in Python on Towards Data Science
- A Conceptual Explanation of Bayesian Hyperparameter Optimization for Machine Learning on Towards Data Science

Practice Questions:

1. Compare grid search and random search methods.
2. What is Bayesian optimization, and how does it improve the search process?

DAY 18

Advanced Machine Learning Topics

Goals:

Dive into more complex ML topics and current trends.

Topics:

- AutoML
- Federated learning
- Explainable AI (XAI)

Resources:

- Automated Machine Learning on Coursera
- Federated Learning: Collaborative Machine Learning without Centralized Training Data on Google AI Blog
- Explainable Artificial Intelligence (XAI): Concepts, Taxonomies, Opportunities and Challenges toward Responsible AI on Information Fusion

Practice Questions:

1. What is AutoML, and how can it automate the machine learning pipeline?
2. Discuss the importance of explainable AI in model transparency.

DAY 19

Integrating Machine Learning into Applications

Goals:

Learn how to deploy machine learning models into production environments.

Topics:

- Model deployment strategies
- Flask for API development
- Real-time prediction systems

Resources:

- Deploying Machine Learning Models: A Complete Guide on Coursera
- Building Machine Learning Web Apps with Flask on FreeCodeCamp

Practice Questions:

1. Describe the steps involved in deploying a machine learning model.
2. How can Flask be used to create APIs for machine learning models?

DAY 20

Review and Project Implementation

Goals:

Review all topics covered and apply knowledge to a comprehensive project.

Topics:

- Comprehensive review of all major topics
- Project: Design and implement a machine learning solution from scratch

Resources:

- Machine Learning Project Checklist by Dr. Jason Brownlee on Machine Learning Mastery
- End-to-End Machine Learning Project Tutorial — Part 1/2 on Towards Data Science

Practice Questions:

1. Outline a project plan that incorporates different machine learning techniques learned.
2. Reflect on the learning process and identify areas for further development.

BOSSCODER ACADEMY

WHY BOSSCODER?

1000+ Alumni placed at Top Product-based companies.

More than **136% hike** for every **2 out of 3** working professional.

Average package of **24LPA.**

The syllabus is most **up-to-date** and the list of problems provided covers **all important topics.**

Lavanya
Meta

Course is very well **structured** and **streamlined** to crack any **MAANG** company

Rahul .
Google

EXPLORE MORE